

Text Analytics: 2nd Assignment

Tsirmpas Dimitris
Drouzas Vasilis

February 11, 2024

Athens University of Economics and Business
MSc in Data Science

Contents

1	Introduction	2
2	POS Tagging	2
2.1	Dataset	2
2.1.1	Acquisition	2
2.1.2	Qualitative Analysis	2
2.1.3	Preprocessing	3
2.2	Baseline Classifier	3
2.3	MLP Classifier	5
2.3.1	Hyper-parameter tuning	5
2.3.2	Training	5
2.3.3	Results	5
3	Sentiment Analysis	6
3.1	Dataset	6
3.1.1	Average Document Length (Before pre-processing)	6
3.1.2	Pre-processing	6
3.1.3	Average Document Length (after pre-processing)	11
3.1.4	Splitting the dataset	11
3.1.5	TF-IDF	11
3.1.6	Feature selection with SVD	11
3.2	Classifiers	12
3.2.1	DummyClassifier	12
3.2.2	Logistic Regression	12
3.3	Our custom MLP classifier	13

1 Introduction

This report will briefly discuss the theoretical background, implementation details and decisions taken for the construction of MLP models for sentiment analysis and POS tagging tasks.

This report and its associated code, analysis and results were conducted by the two authors. Specifically, the sentiment analysis task was performed by Drouzas Vasilis, and the POS-tagging task by Tsirmpas Dimitris. This report was written by both authors.

Note that due to the relative custom code complexity, most of the code used in this section was developed and imported from python source files located in the 'tasks' module. In-depth documentation and implementation details can be found in these files.

2 POS Tagging

POS tagging is a language processing task where words in a given text are assigned specific grammatical categories, such as nouns, verbs, or adjectives. The objective is to analyze sentence structure.

In this section we describe how we can leverage pre-trained word embeddings to create a context-aware MLP classifier.

2.1 Dataset

Acquiring and preprocessing our data with the goal of eventually acquiring a sufficient representation of our text is the most difficult and time-consuming task. We thus split it in distinct phases:

- Original dataset acquisition and parsing
- Qualitative analysis and preprocessing
- Transformation necessary for the NLP task

Each of these distinct steps are individually analyzed below.

2.1.1 Acquisition

We select the [English EWT-UD](#) tree, which is the largest currently supported collection for POS tagging tasks for the English language.

This corpus contains 16622 sentences, 251492 tokens and 254820 syntactic words, as well as 926 types of words that contain both letters and punctuation, such as 's, n't, e-mail, Mr., 's, etc). This is markedly a much higher occurrence than its siblings, and therefore may lead to a slightly more difficult task.

The dataset is made available in `conllu` format, which we parse using the recommended `conllu` python library. We create a dataframe for every word and its corresponding POS tag and link words belonging to the same sentences by a unique sentence ID. The data are already split to training, validation and test sets, thus our own sets correspond to the respective split files.

We are interested in the UPOS (Universal Part of Speech) tags for English words.

2.1.2 Qualitative Analysis

Our training vocabulary is comprised of 16654 words. We include qualitative statistics on the sentences of our dataset in Tables 1 and 2. The splits are explicitly mentioned separately because the splitting was performed by the dataset authors and not by random sampling. We would therefore like to confirm at a glance whether their data are similar.

Set	Mean	Std	Min	25%	50%	75%	Max
Training	16.31	12.4	1	7	14	23	159
Validation	12.56	10.41	1	5	10	17	75
Test	12.08	10.6	1	4	9	17	81

Table 1: Summary and order statistics for the number of words in the sentences of each data split.

Set	Total Word Count	Total Sentence Count
Training	204614	12544
Validation	25152	2001
Test	25096	2077

Table 2: Total text volume of each data split.

2.1.3 Preprocessing

Given the nature of our task we can not implement preprocessing steps such as removing punctuation marks, stopwords or augmenting the dataset. Thus, the only meaningful preprocessing at this stage would be converting the words to lowercase. We believe that the context of each word will carry enough information to distinguish its POS tag regardless of case.

Another issue we need to address before continuing is that of words being part of (depending on) other words for their POS tag. Those would be words such as "don't", "couldn't" or "you're". In the standard UPOS schema these are defined as two or more separate words, where the first is represented by its standard POS tag, and the rest as part of that tag (UPOS tag "PART"). For instance, "don't" would be split into "do" and "n't" with "AUX" and "PART" tags respectively. In our dataset, these words are represented both in the manner described above followed by the full word ("don't") tagged with the pseudo-tag "_". We remove the latter representation from the working dataset.

The general algorithm to calculate the window embeddings on our dataset can be found in Algorithm 1. The algorithm uses a few external functions which are not described here for the sake of brevity. `get_window()` returns the context of the word inside a sentence, including padding where needed, `embedding()` returns the word embedding for a single word and `concatenate` returns a single element from a list of elements. The rest of the functions should be self-explanatory. Note that this algorithm does not represent the actual python implementation.

2.2 Baseline Classifier

We create our own classifier which classifies each token by the majority label associated with it. The classifier is defined as a subclass of sklearn's classifier superclass and thus can seamlessly use it in most sklearn-provided functions such as `classification_report()` and its implementation can be found in the `tasks.models` module.

The results of the classifier can be found in Figure TODO. These results make intuitive sense, since most words in the English language can be classified in a single label, irrespective of context. For example, "is" will always be classified as "AUX", and all punctuation marks will be classified as "PUNCT".

Thus, besides quantitative statistics such as categorical accuracy and f1-score, we should pay close attention to the precision and recall statistics for the more variable POS tags such as "NOUN" or "VERB" in order to properly evaluate our MLP classifier.

Algorithm 1 Window Embedding creation algorithm from raw-text sentences.

Input sentences: a list of sentences

Output tuple(windows, targets): the window embeddings and the POS tag corresponding to the median word of each window

```
1: windows = list()
2: targets = list()
3:
4: for sentence in sentences do
5:     for word in sentence do
6:         window = get_window(word, sentence)
7:         target = get_tag(word)
8:         windows.add(window)
9:         targets.add(target)
10:    end for
11: end for
12:
13: window_embeddings = list()
14: for window in windows do
15:     word_embeddings = list()
16:     for word in window do
17:         if word is PAD_TOKEN then
18:             word_embeddings.add(zeros(embedding_size))
19:         else
20:             word_embeddings.add(embedding(word))
21:         end if
22:     end for
23:     window_embedding = concatenate(word_embeddings)
24:     window_embeddings.add(window_embedding)
25: end for
26:
27: return window_embeddings, one_hot(targets)
```

2.3 MLP Classifier

2.3.1 Hyper-parameter tuning

We use the `keras_tuner` library to automatically perform random search over various hyper-parameters of our model.

The parameter search consists of:

- The depth of the model (the number of layers)
- The height of the model (the number of parameters by layer)
- The learning rate

The parameter search does NOT consist of:

- Dropout rate, since dropout rarely changes the final result of a neural network, but rather tunes the trade-off between training time and overfit avoidance
- Activation functions, since they rarely significantly influence the model’s performance

With this scheme we hope to maximize the area and granularity of our search to the hyper-parameters that are most likely to significantly influence the final results.

We implement early stopping and set a maximum iteration limit of 70. We assume that if a model needs to go over that limit, it may be computationally inefficient, and thus less desirable compared to a slightly worse, but much more efficient model. Additionally, we use a relatively large batch size to improve training times since this operation is very computationally heavy. We don’t yet aim to create the best classifier, so slightly suboptimal weights are not a problem for the purposes of the hyperparameter search. We use a relatively very large batch size to improve training times since this operation is very computationally heavy. We don’t yet aim to create the best classifier, so slightly suboptimal weights are not a problem for the purposes of the hyperparameter search.

2.3.2 Training

We now re-train our model with a much smaller batch size and keep track of the training history and best weights by validation loss.

Unfortunately, the different batch size means we can not rely on the hyper parameter search to get an estimation of training epochs. Thus, we rely on early-stopping on the validation data to ensure our model does not overfit as a result of training time.

We use the categorical accuracy stopping criterion instead of loss. This may lead to situations where validation loss increases, but so does accuracy [1]. This represents a trade-off between our model being more confidently incorrect about already-misclassified instances, but better at edge cases where the classification is more ambiguous. We previously discussed how the strength of a context-aware classifier lies in these kinds of distinctions, which justifies our choice of favoring correct edge-case classifications in the expense of more confidently incorrect misclassifications.

This phenomenon is demonstrated in Figure 1.

2.3.3 Results

The results of our MLP classifier compared to the baseline model mentioned above can be found in Tables 3, 4 and 5. We include precision, recall and F1 scores for each individual tag, as well as their macro average denoted by the "MACRO" tag in the tables.

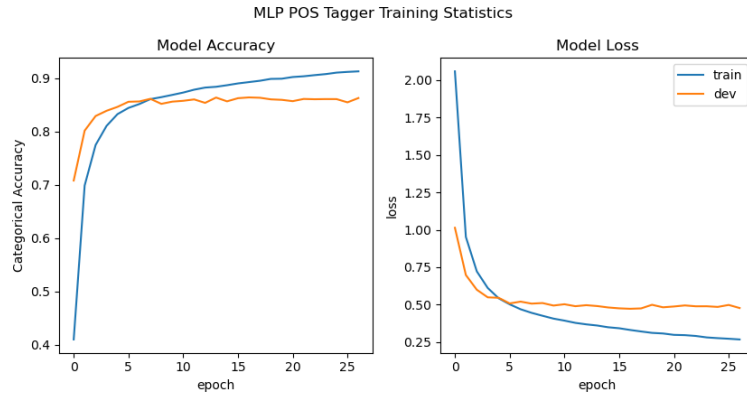


Figure 1: Loss and accuracy on the training and validation sets depending on the number of epochs.

```

---Before preprocessing---
Average Document Length (in words): 746.3405
Average Document Length (in characters): 3893.002

```

Figure 2: Average Document length in words and characters (before the pre-processing).

We note an increase in all metrics for our MLP classifier, especially in tags such as "SCONJ" (subordinating conjunction), and VERB and NOUN (which we hypothesized at the beginning of the report). The only notable exception is the "X" tag (other) which is attributed to unintelligible material, foreign words and word fragments. This is an acceptable drawback, since these are easily caught by preprocessing or weak classifiers anyway.

3 Sentiment Analysis

Sentiment analysis, also known as opinion mining, is the process of analyzing text to determine the sentiment or emotional tone expressed within it. The goal of sentiment analysis is to understand the attitudes, opinions, and emotions conveyed by the text.

3.1 Dataset

Here we will be working with the Cornell Movie Review dataset, which consists of 2000 movie reviews, split equally in 1000 positive and 1000 negative ones. The goal here will be to develop classifiers that will effectively understand whether a review is a positive or negative one, based on the data it has been trained on. We begin by taking a brief look into our dataset.

3.1.1 Average Document Length (Before pre-processing)

The average document length in words and characters before the processing can be checked in Figure 2.

3.1.2 Pre-processing

For demonstration reasons, we start by printing the 20 most frequent words in the text, in Figure 3.

Table 3: Results on the training dataset.

model	tag	precision	recall	f1
Baseline	ADJ	1.000	0.886	0.940
Baseline	ADP	1.000	0.826	0.905
Baseline	ADV	1.000	0.841	0.913
Baseline	AUX	1.000	0.924	0.960
Baseline	CCONJ	1.000	0.994	0.997
Baseline	DET	1.000	0.958	0.979
Baseline	INTJ	1.000	0.917	0.957
Baseline	NOUN	1.000	0.908	0.952
Baseline	NUM	1.000	0.985	0.992
Baseline	PART	1.000	0.884	0.938
Baseline	PRON	1.000	0.959	0.979
Baseline	PROPN	1.000	0.894	0.944
Baseline	PUNCT	1.000	0.997	0.999
Baseline	SCONJ	1.000	0.455	0.626
Baseline	SYM	1.000	0.835	0.910
Baseline	VERB	1.000	0.899	0.947
Baseline	X	1.000	0.676	0.806
Baseline	MACRO	0.863	0.873	0.862
MLP	ADJ	1.000	0.919	0.958
MLP	ADP	1.000	0.869	0.930
MLP	ADV	1.000	0.904	0.949
MLP	AUX	1.000	0.982	0.991
MLP	CCONJ	1.000	0.995	0.998
MLP	DET	1.000	0.986	0.993
MLP	INTJ	1.000	0.824	0.903
MLP	NOUN	1.000	0.945	0.972
MLP	NUM	1.000	0.969	0.984
MLP	PART	1.000	0.990	0.995
MLP	PRON	1.000	0.938	0.968
MLP	PROPN	1.000	0.842	0.914
MLP	PUNCT	1.000	0.997	0.999
MLP	SCONJ	1.000	0.668	0.801
MLP	SYM	1.000	0.969	0.984
MLP	VERB	1.000	0.889	0.941
MLP	X	1.000	0.500	0.667
MLP	MACRO	0.881	0.893	0.885

Table 4: Results on the validation dataset.

model	tag	precision	recall	f1
Baseline	ADJ	1.000	0.751	0.858
Baseline	ADP	1.000	0.838	0.912
Baseline	ADV	1.000	0.748	0.856
Baseline	AUX	1.000	0.922	0.960
Baseline	CCONJ	1.000	0.990	0.995
Baseline	DET	1.000	0.951	0.975
Baseline	INTJ	1.000	0.690	0.816
Baseline	NOUN	1.000	0.901	0.948
Baseline	NUM	1.000	0.704	0.827
Baseline	PART	1.000	0.865	0.928
Baseline	PRON	1.000	0.964	0.982
Baseline	PROPN	1.000	0.395	0.566
Baseline	PUNCT	1.000	0.989	0.994
Baseline	SCONJ	1.000	0.453	0.624
Baseline	SYM	1.000	0.812	0.896
Baseline	VERB	1.000	0.746	0.855
Baseline	X	1.000	0.020	0.040
Baseline	MACRO	0.795	0.749	0.759
MLP	ADJ	1.000	0.812	0.896
MLP	ADP	1.000	0.847	0.917
MLP	ADV	1.000	0.769	0.870
MLP	AUX	1.000	0.947	0.973
MLP	CCONJ	1.000	0.975	0.988
MLP	DET	1.000	0.971	0.985
MLP	INTJ	1.000	0.598	0.748
MLP	NOUN	1.000	0.865	0.928
MLP	NUM	1.000	0.872	0.931
MLP	PART	1.000	0.959	0.979
MLP	PRON	1.000	0.916	0.956
MLP	PROPN	1.000	0.630	0.773
MLP	PUNCT	1.000	0.993	0.996
MLP	SCONJ	1.000	0.594	0.746
MLP	SYM	1.000	0.884	0.938
MLP	VERB	1.000	0.808	0.894
MLP	X	1.000	0.173	0.296
MLP	MACRO	0.788	0.801	0.791

Table 5: Results on the test dataset.

model	tag	precision	recall	f1
Baseline	ADJ	1.000	0.754	0.860
Baseline	ADP	1.000	0.826	0.905
Baseline	ADV	1.000	0.810	0.895
Baseline	AUX	1.000	0.910	0.953
Baseline	CCONJ	1.000	0.997	0.998
Baseline	DET	1.000	0.956	0.978
Baseline	INTJ	1.000	0.758	0.862
Baseline	NOUN	1.000	0.897	0.946
Baseline	NUM	1.000	0.584	0.737
Baseline	PART	1.000	0.887	0.940
Baseline	PRON	1.000	0.971	0.985
Baseline	PROPN	1.000	0.430	0.602
Baseline	PUNCT	1.000	0.990	0.995
Baseline	SCONJ	1.000	0.438	0.609
Baseline	SYM	1.000	0.826	0.905
Baseline	VERB	1.000	0.766	0.867
Baseline	X	0.000	0.000	0.000
Baseline	MACRO	0.779	0.753	0.755
MLP	ADJ	1.000	0.810	0.895
MLP	ADP	1.000	0.843	0.915
MLP	ADV	1.000	0.817	0.899
MLP	AUX	1.000	0.959	0.979
MLP	CCONJ	1.000	0.990	0.995
MLP	DET	1.000	0.974	0.987
MLP	INTJ	1.000	0.537	0.698
MLP	NOUN	1.000	0.864	0.927
MLP	NUM	1.000	0.836	0.911
MLP	PART	1.000	0.959	0.979
MLP	PRON	1.000	0.913	0.955
MLP	PROPN	1.000	0.617	0.763
MLP	PUNCT	1.000	0.988	0.994
MLP	SCONJ	1.000	0.583	0.736
MLP	SYM	1.000	0.917	0.957
MLP	VERB	1.000	0.828	0.906
MLP	X	1.000	0.226	0.368
MLP	MACRO	0.800	0.803	0.795

```
,: 77717 occurrences
the: 76276 occurrences
.: 65876 occurrences
a: 37995 occurrences
and: 35404 occurrences
of: 33972 occurrences
to: 31772 occurrences
is: 26054 occurrences
in: 21611 occurrences
's: 18128 occurrences
``: 17625 occurrences
it: 16059 occurrences
that: 15912 occurrences
): 11781 occurrences
(: 11664 occurrences
as: 11349 occurrences
with: 10782 occurrences
for: 9918 occurrences
this: 9573 occurrences
his: 9569 occurrences
```

Figure 3: Average Document length in words and characters (before the pre-processing).

```
['not',
 "don't",
 "aren't",
 "couldn't",
 "didn't",
 "doesn't",
 "hadn't",
 "hasn't",
 "shouldn't",
 "haven't",
 "wasn't",
 "weren't",
 "isn't",
 'doesn']
```

Figure 4: Average Document length in words and characters (before the pre-processing).

Most of these words are actually stop words. As in most text classification problems, we would typically need to remove the stop words of the text.

The english stopwords is a package of 179 words that in general, would not help in a sentiment analysis problem. But, since they include terms that are negative, removing them could prove harmful for our case, since we are dealing with a sentiment analysis problem.

e.g. imagine the phrase "I didn't like the film" to end up "like film". Disastrous, right?

So, the plan is to remove all the stop words that include negative meaning before the preprocessing. The stop words that we decided to keep in the text are shown in Figure 4 :

Moving on to the pre-processing task, the steps performed are the following:

- Combination to a single document.
- Conversion to lowercase.
- Lemmatization and stop words extraction.
- Punctuation removal.

```
---After preprocessing---
Average Document Length (in words): 746.3405
Average Document Length (in characters): 3893.002
```

Figure 5: Average Document length in words and characters (before the pre-processing).

Set	Total Word Count	Total Document Count
Training	36624	1400
Validation	16948	300
Test	16780	300

Table 6: Total text volume of each data split.

- Number removal.
- Single characters removal.
- Converting multiple spaces to single ones.

3.1.3 Average Document Length (after pre-processing)

The final average document length is shown in Figure 5.

3.1.4 Splitting the dataset

We decided to split the (processed) dataset into the training set (70%), development set (15%) and test set (15%). The sizes of each set are shown in Table 6.

3.1.5 TF-IDF

We used the unigram and bi-gram TF-IDF features, defining the maximum number of features to 5000. The shapes of the data are shown in Figure 6.

3.1.6 Feature selection with SVD

We performed dimensionality reduction with the TruncatedSVD() method, reducing the number of features from 5000 to 500. The new shapes can be found in Figure 7.

```
Shape (training data): (1400, 5000)
Shape (development data): (300, 5000)
Shape (test data): (300, 5000)
```

Figure 6: The shapes of the TF-IDF vectors.

```

Shape (training data) after SVD: (1400, 500)
Shape (development data) after SVD: (300, 500)
Shape (test data) after SVD: (300, 500)

```

Figure 7: The new shapes after doing SVD analysis.

Classification Report on Development Set:				
	precision	recall	f1-score	support
neg	0.00	0.00	0.00	157
pos	0.48	1.00	0.65	143
accuracy			0.48	300
macro avg	0.24	0.50	0.32	300
weighted avg	0.23	0.48	0.31	300

Classification Report on Training Set:				
	precision	recall	f1-score	support
neg	0.00	0.00	0.00	690
pos	0.51	1.00	0.67	710
accuracy			0.51	1400
macro avg	0.25	0.50	0.34	1400
weighted avg	0.26	0.51	0.34	1400

Classification Report on Test Set:				
	precision	recall	f1-score	support
neg	0.00	0.00	0.00	153
pos	0.49	1.00	0.66	147
accuracy			0.49	300
macro avg	0.24	0.50	0.33	300
weighted avg	0.24	0.49	0.32	300

Figure 8: Classification results of DummyClassifier for training, test and validation sets.

3.2 Classifiers

3.2.1 DummyClassifier

DummyClassifier makes predictions that ignore the input features. This classifier serves as a simple baseline to compare against other more complex classifiers. The strategy to generate predictions was set to 'most_frequent', meaning that the predict method always returns the most frequent class label. The results of this classifier are demonstrated in Figure 8.

As expected, the results are poor since the decision of the classifier depends exclusively only the majority class.

3.2.2 Logistic Regression

Logistic Regression is a statistical method used for binary classification tasks, where the output variable takes only two possible outcomes. Before applying Logistic Regression, we will perform a grid search to find the optimal parameters to run the classifier. The parameters we tried are the following:

- Solver: We tested 'liblinear' and 'saga' solvers
- Penalty: We tested 'l1', 'l2', 'elasticnet' regularization penalties

```

0s [31] Classification Report on Training Set:
      precision    recall  f1-score   support

      neg      0.94      0.92      0.93      690
      pos      0.92      0.94      0.93      710

   accuracy      0.93      0.93      0.93      1400
  macro avg      0.93      0.93      0.93      1400
 weighted avg      0.93      0.93      0.93      1400

-----

Classification Report on Development Set:
      precision    recall  f1-score   support

      neg      0.90      0.82      0.85      157
      pos      0.82      0.90      0.85      143

   accuracy      0.86      0.86      0.85      300
  macro avg      0.86      0.86      0.85      300
 weighted avg      0.86      0.85      0.85      300

-----

Classification Report on Test Set:
      precision    recall  f1-score   support

      neg      0.88      0.81      0.84      153
      pos      0.82      0.88      0.85      147

   accuracy      0.85      0.85      0.85      300
  macro avg      0.85      0.85      0.85      300
 weighted avg      0.85      0.85      0.85      300

```

Figure 9: Metrics of the Logistic Regression on the training, test and development sets.

- C: We tested values of 0.001, 0.01, 0.1, 1 and 10 (inverse of regularization strength)

The best hyperparameters were the following: C= 10, penalty= 'l1', solver = 'liblinear'.

Now, it is time to fit the Logistic Regression using these parameters. The results we got are shown in Figure 9 .

3.3 Our custom MLP classifier

First of all, we define the `y_train_1_hot` and `y_dev_1_hot` vectors using the `LabelBinarizer` and applying `fit_transform()` and `transform()` to the training and development 1-hot vectors respectively. Now, it's time to define our MLP model. We followed the same method as before, i.e. we used the `keras_tuner` library to perform random search with parameters the depth of the model, the height and the learning rate. Early stopping is implemented , with main criterion the validation loss. The number of epochs was defined to 100. We experimented with a variety of different hyperparameter combinations. The first experiments examined the quantity of the hidden layers, keeping other parameters constant. The results tended to optimize using only one hidden layer, meaning that the problem can be solved with a small network. After defining the number of hidden layers to one, we examined other hyperparameter combinations:

Model	1	2	3	4	5	6*	7*
#Hidden layers	1	1	1	1	1	1	1
Batch size	256	256	256	128	512	256	256
Min_hid_size	64	32	16	32	32	32	32
Max_hid_size	1024	512	206	512	512	512	512

- Model 1: Good loss, mediocre accuracy.



Figure 10: MLP accuracy as a function of epochs.

- Model 2: Optimal loss & accuracy.
- Model 3: Worse results vs Model (2) in terms of loss.
- Model 4: Data tend to overfit.
- Model 5: Worse results compared to Model (2) in both accuracy,loss.
- Model 6 (*removing the subsequent layer): Similar results to model (2).
- Model 7(*Learning rate set to 1e-5): Worse results compared to model(2).

The model we chose to keep is Model 2, which seems to be the optimal one after testing the hyperparameters. The results we gain are shown in Figures 10, 11.

At a first glance, the curves showing the model accuracy do seem a bit weird. Specifically, we would more like expect the training accuracy to start at a lower rate and increase over time. But what we see is that it actually starts already pretty high (96.71%) and ends up slightly higher (98.36%). It means our model very accurately predicts the training data even from the beginning.

As far as the validation data are concerned, we see that no significant improvements are made while the number of epochs increases.

In general, the binary accuracies for both training and validation sets are relatively high and show similar trends of improvement over epochs. Additionally, there is no significant gap between the training and validation accuracies. This suggests that the model is not overfitting to the training data.

Our prediction is that the data are already quite simple to distinguish, so the parameters are already too many to learn to distinguish. This is why the development accuracy does not show actually an improvement over time. Now regarding the training/development losses, here we see that both training and validation losses decrease together, which suggests that the model is generalizing well.




Figure 11: MLP loss as a function of epochs.

Class neg :	(Training)	(Development)	(Test)
Precision	0.992733	0.912162	0.848684
Recall	0.989855	0.859873	0.843137
F1-score	0.991292	0.885246	0.845902
PR AUC	0.999691	0.951337	0.922453
<hr/>			
Class pos :	(Training)	(Development)	(Test)
Precision	0.990169	0.855263	0.837838
Recall	0.992958	0.909091	0.843537
F1-score	0.991561	0.881356	0.840678
PR AUC	0.999691	0.951337	0.922453

Figure 12: Metrics for the MLP classifier for both classes for the training, development and test sets.


```

Macro-averaged Scores for Training Subset:
=====
Macro-averaged Precision: 0.991451
Macro-averaged Recall: 0.991406
Macro-averaged F1-score: 0.991426
Macro-averaged PR AUC: 0.999691

Macro-averaged Scores for Development Subset:
=====
Macro-averaged Precision: 0.883713
Macro-averaged Recall: 0.884482
Macro-averaged F1-score: 0.883301
Macro-averaged PR AUC: 0.951337

Macro-averaged Scores for Test Subset:
=====
Macro-averaged Precision: 0.843261
Macro-averaged Recall: 0.843337
Macro-averaged F1-score: 0.843290
Macro-averaged PR AUC: 0.922453

```

Figure 13: Macro-Metrics for the MLP classifier for both classes for the training, development and test sets.

Next, we will provide the metrics (Precision, Recall, F1 score and the AUC scores) for training, development and test subsets in Figure 12.

Finally, let's take a look at the Macro-averaged metrics (averaging the corresponding scores of the previous bullet over the classes) for the training, development and test subsets, as shown in Figure 13.

References

- [1] Soltius (<https://stats.stackexchange.com/users/201218/soltius>). *How is it possible that validation loss is increasing while validation accuracy is increasing as well*. Cross Validated. URL: <https://stats.stackexchange.com/q/341054>. (version: 2023-03-28). eprint: <https://stats.stackexchange.com/q/341054>. URL: <https://stats.stackexchange.com/q/341054>.